

# American Sign Language (ASL) Detection – Project Report

## 1. Introduction

This project focuses on building a computer vision system that can recognize **American Sign Language (ASL)** hand gestures. The goal was to classify 29 different signs (A–Z, SPACE, DELETE, NOTHING) using a deep learning model and deploy it through a simple Streamlit web app.

## 2. Dataset

The dataset used is the **ASL Alphabet Dataset from Kaggle**, which contains:

- 29 classes
- Around **3,000 training images per class**
- A separate folder for test images

The images cover a wide variety of hand positions, backgrounds, and lighting conditions, which helps the model learn more robust features.

## 3. Methodology

### *a. Preprocessing*

- All images were resized to **64×64** for faster CPU training.
- Pixel values were normalized to the range **0–1**.
- The dataset was split using TensorFlow’s validation\_split feature (80% training, 20% validation).

### *b. Model Selection*

To keep the project lightweight and workable on a normal CPU, I used **MobileNetV2** (pre-trained on ImageNet) with frozen base layers.

A custom classification head was added on top with:

- Global Average Pooling

- Dropout layer
- Dense output layer (29 units)

This allowed efficient transfer learning without needing heavy GPU resources.

### c. *Training*

The model was trained for **5 epochs** due to system limitations.

Optimizer: **Adam**

Loss: **Categorical Crossentropy**

## 4. Results

- **Training Accuracy:** ~91%
- **Validation Accuracy:** ~64%

There is a noticeable gap between training and validation accuracy. This indicates **early overfitting**, which is expected because the model was trained for a limited number of epochs on a CPU.

With more time or GPU training, the validation accuracy is expected to improve significantly (typically 85–95% for this dataset).

## 5. Overfitting Observation

During training, the model performed extremely well on the training images but struggled a bit with unseen data.

This happens because:

- The dataset is large (around 87,000 images total).
- The model had limited time to learn generalizable features.
- Training was done on a standard laptop CPU.

Despite this, the model still performs decently as a **prototype** and correctly predicts many ASL test images.

## 6. Deployment

A simple **Streamlit app** (app.py) was built where users can upload an image, and the system displays the predicted ASL letter.

This makes the project interactive and easy to demonstrate.

## 7. Conclusion

This project successfully demonstrates a complete pipeline for ASL sign detection—from dataset preprocessing and model training to deployment through a web interface.

Even with restricted hardware, the model achieved promising results and proves the feasibility of real-time ASL recognition.

With further training on GPU and fine-tuning more layers of MobileNetV2, performance can be improved significantly.

## 8. Future Improvements

- Train for more epochs using GPU
- Add stronger data augmentation
- Fine-tune deeper layers of MobileNetV2
- Implement real-time webcam-based detection